Slide 1

Welcome to Developing and Deploying Intelligent Chat Bots

Instructor: Micheleen Harris
Data Scientist @ Microsoft

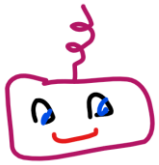

Slide 2

Technologies covered

- Microsoft Cognitive Services
- Microsoft Bot Framework

Slide 3

What you'll know at the end of Cognitive Services section

1. Understand the types of intelligent services offered by Microsoft Cognitive Services
2. Be able to talk about and use the Language Understanding and Intelligent Service API
3. Be able to talk about and use the Computer Vision API
4. Be able to talk about and use the Bing News Search API

Slide 4

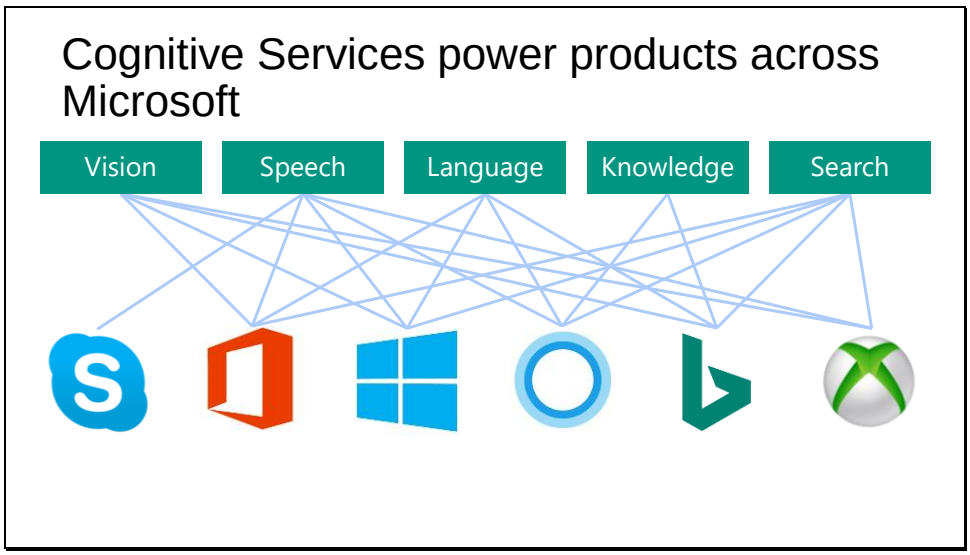

Cognitive Services are coming out of as well as being consumed by groups across Microsoft including our products as well as Microsoft Research and engineering

e.g. Text Analytics under Language has come from engineers in Azure – Microsoft's cloud offering
e.g. the Knowledge Exploration Service for building custom knowledge bases has come out of Microsoft Research
e.g. Speech to text has come out of Bing

Slide 5

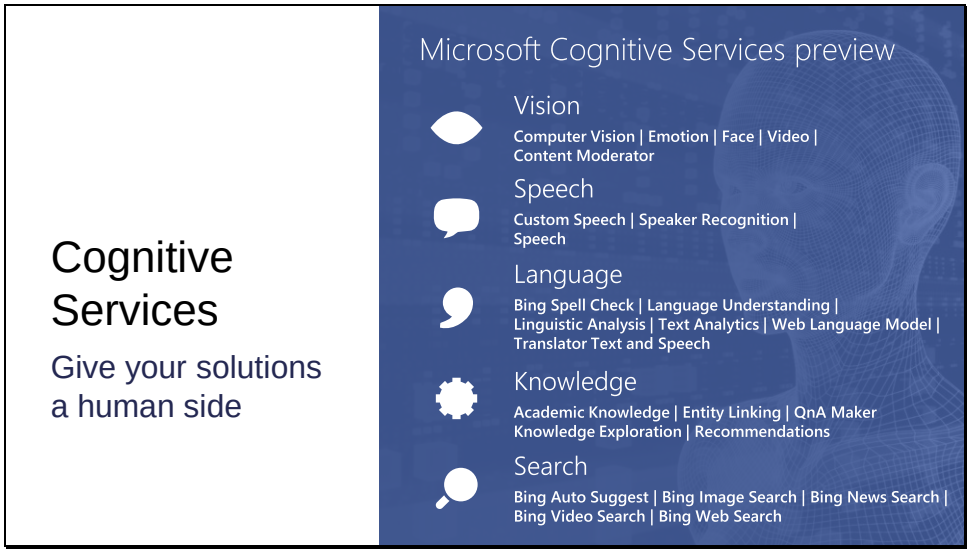

At Microsoft, we've been offering APIs for a very long time across the company. In delivering Microsoft Cognitive Services API, we started with 4 at /build (2015); added 7 more December 2015, and today we have 24 APIs in our collection.

Cognitive Services are available individually or as a part of the Cortana Intelligence Suite, formerly known as Cortana Analytics, which provides a comprehensive collection of services powered by cutting-edge research into machine learning, perception, analytics and social bots.

These APIs are powered by Microsoft Azure.

Developers and businesses can use this suite of services and tools to create apps that learn about our world and interact with people and customers in personalized, intelligent ways.

Some great scenarios for all of the cognitive services can be found on this wiki: https://github.com/Azure/bot-education/wiki/Cognitive-Services-Scenarios

Slide 6

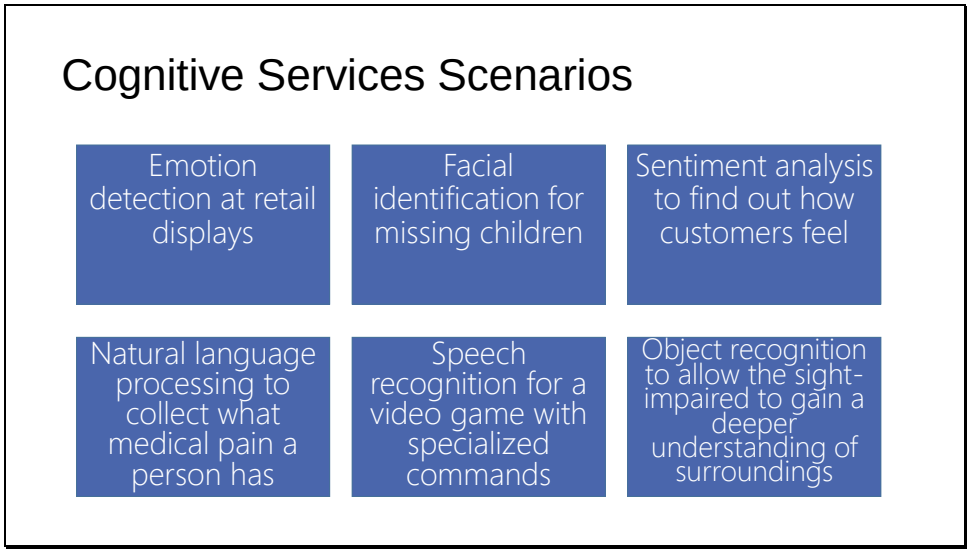
Cognitive Services Scenarios
Emotion detection at retail displays
Facial identification for missing children
Sentiment analysis to find out how customers feel
Natural language processing to collect what medical pain a person has
Speech recognition for a video game with specialized commands
Object recognition to allow the sight-impaired to gain a deeper understanding of surroundings

Slide 7

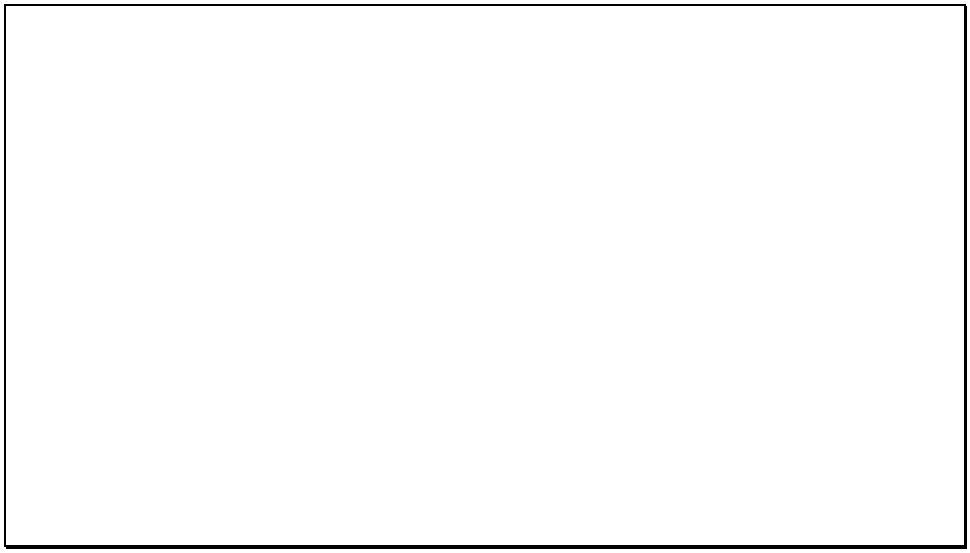

Slide 8

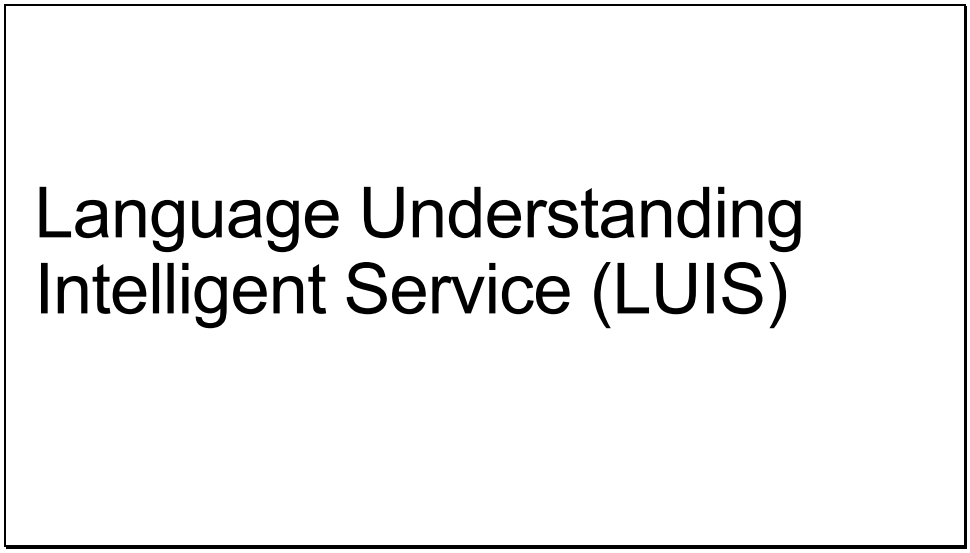

https://www.microsoft.com/cognitive-services/en-us/language-understanding-intelligent-service-luis

Slide 9

How to talk about it
LUIS Concepts

Intent – aim or goal

Entities – a type or “notion” of person, place or thing

Utterances – the phrase we might use that is added training data

Example:
intent – find news on topic and possibly share with another person
entities - We'd like to be able to say what kind of news we are interested in, and also, for sharing, to say who we'd like to share a story with. In order to capture the notion of a news topic, and a recipient for sharing, let's create two entity types: "Topic" and "Recipient".
Utterence – “Get me the news on electric cars and share with Merinda”

We currently support (3/28/2017):

30 entities per application
80 intents per application

Over 2000 utterances per application.

Slide 10

How to understand it

LUIS examples

Understand what your users are saying

- Determines intent
- Detects entities

Learns over time

- Active and interactive learning

Options to use pre-built, world class models from Bing and Cortana

Slide 11

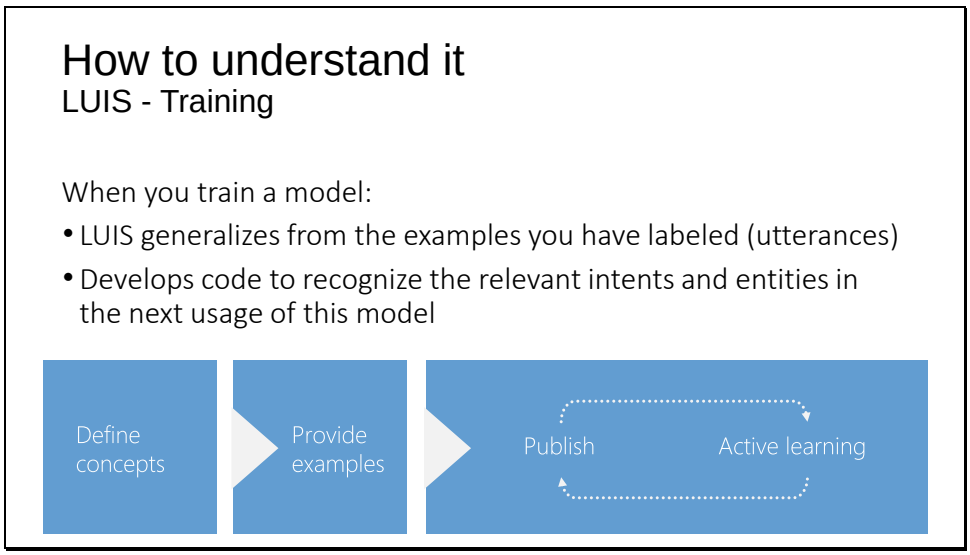

Internally, LUIS uses logistic regression classifiers to determine intents, and conditional random fields (CRFs) to determine the entities. The training process results in optimized classifiers and CRFs, referred to as models, that LUIS can use in the future.

Slide 12

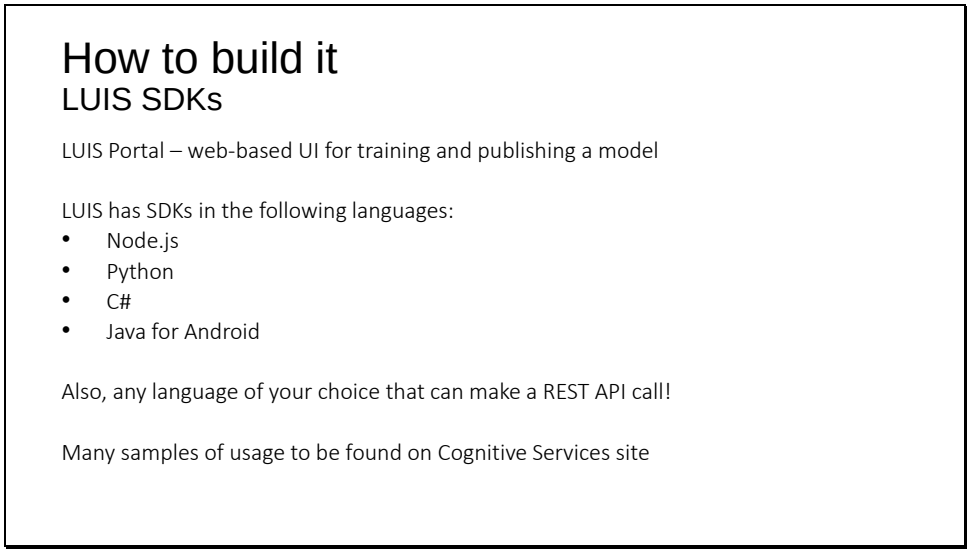

Note on Utterances
If you have unlabeled utterances that your application should handle, they will be available when you edit the application under the "Search" and "Suggest" tabs.

Can link intents to actions in your custom code (consuming the API)

Exporting LUIS app
Can download your work into a JSON file. This lets you share you application with other developers, or check your LUIS application into your version control.

Pre-Built:
LUIS also provides access to pre-built LUIS applications that use many of the same models found in Microsoft Cortana.

SDKs
https://github.com/Microsoft/Cognitive-LUIS-Node.js
https://github.com/Microsoft/Cognitive-LUIS-Python
https://github.com/Microsoft/Cognitive-LUIS-Windows
https://github.com/Microsoft/Cognitive-LUIS-Android

SDK/Samples site:

https://www.microsoft.com/cognitive-services/en-us/SDK-Sample

Slide 13

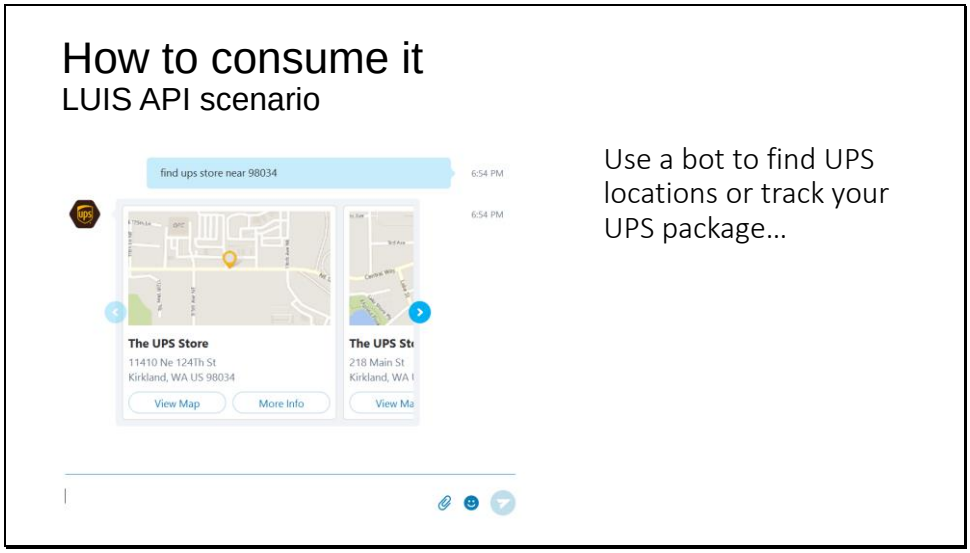

UPS Bot – uses natural language to help a user find what they need along with other APIs
https://bots.botframework.com/bot?id=UPSBot

More
Bot: (Node.js) https://github.com/Microsoft/BotBuilder-Samples/tree/master/Node/intelligence-LUIS
Bot: (C#) https://github.com/Microsoft/BotBuilder-Samples/tree/master/CSharp/intelligence-LUIS
Cognitive Services site demo: https://www.microsoft.com/cognitive-services/en-us/language-understanding-intelligent-service-luis
See LUIS help at https://luis.ai for more

Slide 14

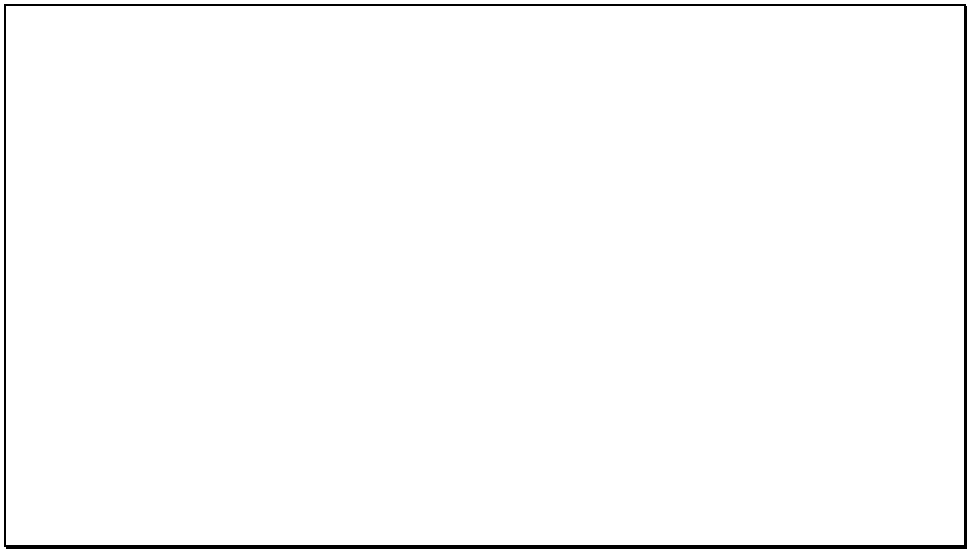

Slide 15

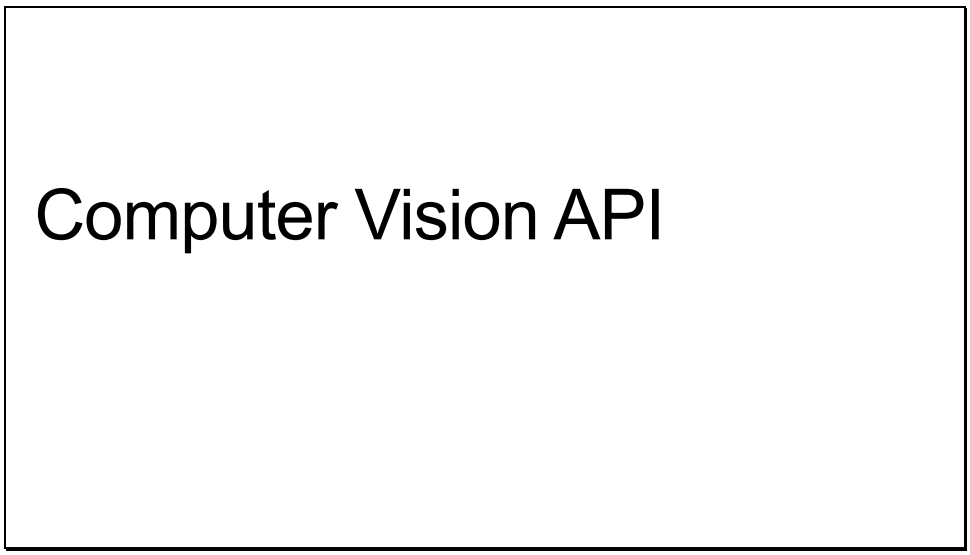

https://www.microsoft.com/cognitive-services/en-us/computer-vision-api

Slide 16

How to talk about it
CV API modes

Modes for analysis:

- Analyze – extraction of visual features (such as tags, colors, faces)
- Describe – natural language description
- OCR – word recognition
- Domain specific – celebrity recognition only currently
- Thumbnail – user desired size with regions of interest

https://www.microsoft.com/cognitive-services/en-us/Computer-Vision-API/documentation -> details in API Reference

Slide 17

How to understand it

Image analysis example

Description	{ "type": 0, "captions": [{ "text": "person that is standing in the snow", "confidence": 0.4835140399902179 }] }
Clip Art Type	0 Non-clipart
Line Drawing Type	0 Non-LineDrawing
Black & White Image	False
Categories	[{ "name": "people_", "score": 0.9375 }]
Faces	[{ "age": 25, "gender": "Female", "faceRectangle": { "width": 288, "height": 288, "left": 874, "top": 557 } }]

Image analyzed with the demo found at https://www.microsoft.com/cognitive-services/en-us/computer-vision-api

Slide 18

How to understand it

Machine Learning models for image captions with CV

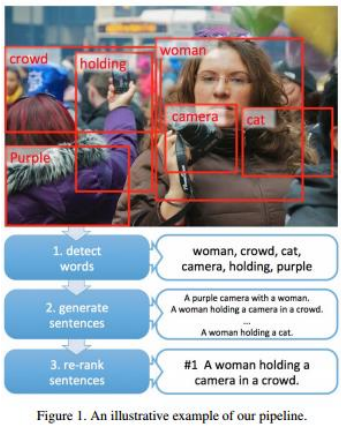

Figure 1. An illustrative example of our pipeline.

Key notes:

- sub-regions are used in learned detector
- Training of language model captures commonsense knowledge
- High likelihood sentences are re-ranked by linear weighting

Fang H. *et al.* From Captions to Visual Concepts and Back. 2016

See paper here: https://www.microsoft.com/en-us/research/wp-content/uploads/2016/02/CVPR15_0866.pdf

Slide 19

How to build it
CV tools and samples

CV has SDKs for the following:

- Swift
- Android (Java)
- Windows (C#)

And samples:

- Python (Jupyter notebook)
- C# (image captioning bot)
- Node.js (image captioning bot)
- See4Me (Face, Emotion, CV – cross-platform app)

Also, any language of your choice that can make a REST API call!

SDKs:
https://github.com/DanilaVladi/Microsoft-Cognitive-Services-Swift-SDK
https://github.com/Microsoft/Cognitive-vision-android
https://github.com/Microsoft/Cognitive-vision-windows

Notebook:
https://github.com/Microsoft/Cognitive-vision-python

Bots:
https://github.com/Microsoft/BotBuilder-Samples/tree/master/Node/intelligence-ImageCaption
https://github.com/Microsoft/BotBuilder-Samples/tree/master/CSharp/intelligence-ImageCaption

See4Me:
https://github.com/DotNetToscana/See4Me

Demos also on CV landing page at: https://www.microsoft.com/cognitive-services/en-us/computer-vision-api

Slide 20

How to consume it

CV scenario

Seeing AI: In collaboration with Microsoft Research, Pivothead is bringing "vision" to the visually-impaired.

By combining APIs from Microsoft Cognitive Services with the imaging performance and power of Pivothead SMART, a person who is visually-impaired can better understand who and what is going on around them.

http://www.pivothead.com/seeingai/

Slide 21

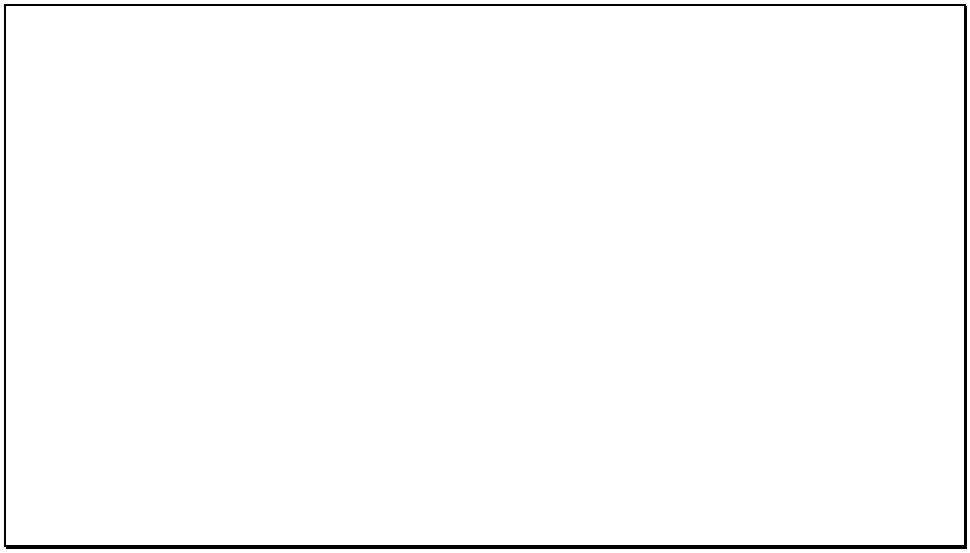

Slide 22

https://www.microsoft.com/cognitive-services/en-us/bing-news-search-api

Slide 23

How to talk about it

News Search API modes

- Search – news articles from a given query
- Category – news articles of a given preset category (e.g. science and technology, world, etc.)
- Trending topics – trending news articles

Slide 24

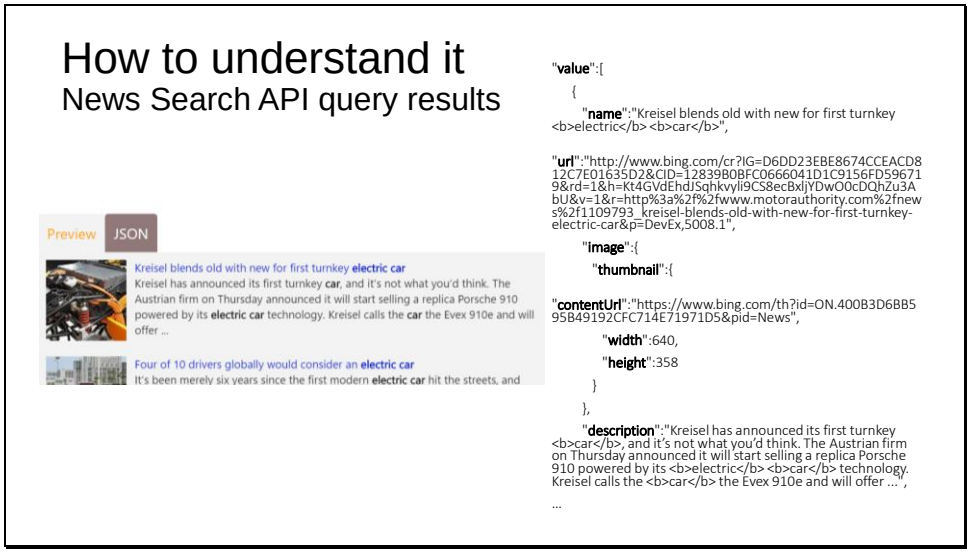

Search the web for news articles. Results include details like authoritative image of the news article, related news and categories, provider info, article URL, and date added.

Slide 25

How to build it

News Search API call

Call the News Search REST API endpoint in a few lines of code…

```
        // Build the url we'll be calling to get top news
        var url = "https://api.cognitive.microsoft.com/bing/v5.0/news/?"
            + "category=" + results.response.entity + "&count=10&mkt=en-US&originalImg=true";
        // Build options for the request
        var options = {
            uri: url,
            headers: {
                'Ocp-Apim-Subscription-Key': BINGSEARCHKEY
            },
            json: true // Returns the response in json
        }
//Make the call
rp(options).then(function (body){
```

Slide 26

How to consume it

News Search scenario

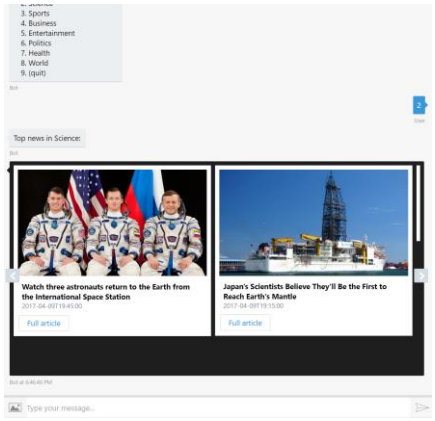

Use a bot to find the news around the world…

Bot: (C#) https://docs.botframework.com/en-us/bot-intelligence/search/#example-trending-news-bot
Bot: (Node.js) https://github.com/alyssaong1/HOL-NewsBot
Cognitive Services Bing News site demo: https://www.microsoft.com/cognitive-services/en-us/bing-news-search-api
Intelligent kiosk application: https://github.com/Microsoft/Cognitive-Samples-IntelligentKiosk
See the API Reference for more details on REST API calls

Slide 27

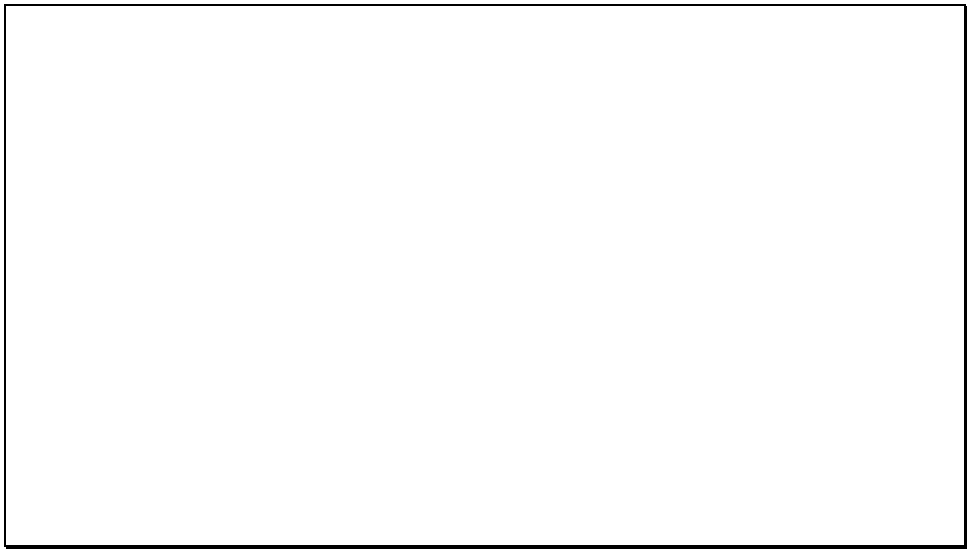

Slide 28

Developer Resources

Pricing

https://www.microsoft.com/cognitive-services Click on Pricing

Documentation

https://www.microsoft.com/cognitive-services Click on Docs + Help

Client SDKs and Example Code

https://www.microsoft.com/cognitive-services/en-us/SDK-Sample

Community

https://stackoverflow.com/questions/tagged/microsoft-cognitive
https://cognitive.uservoice.com/

Slide 29

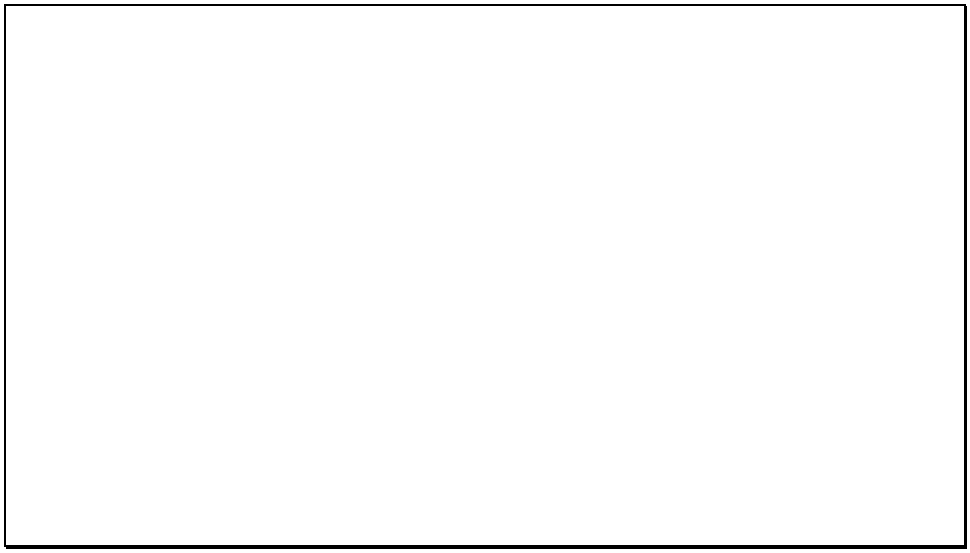

Slide 30

Here enters the Microsoft Bot Framework

Slide 31

What you'll know at the end of the Bot Framework section

1. Gain an overview of the tools and resources needed to build a bot with Microsoft's Bot Framework
2. Gain the knowledge on how to design a great bot experience
3. Understand what components make up the Microsoft Bot Framework
4. See how data works in the Bot Framework
5. Understand how to deploy a bot as an app to Azure, Microsoft's cloud

Slide 32

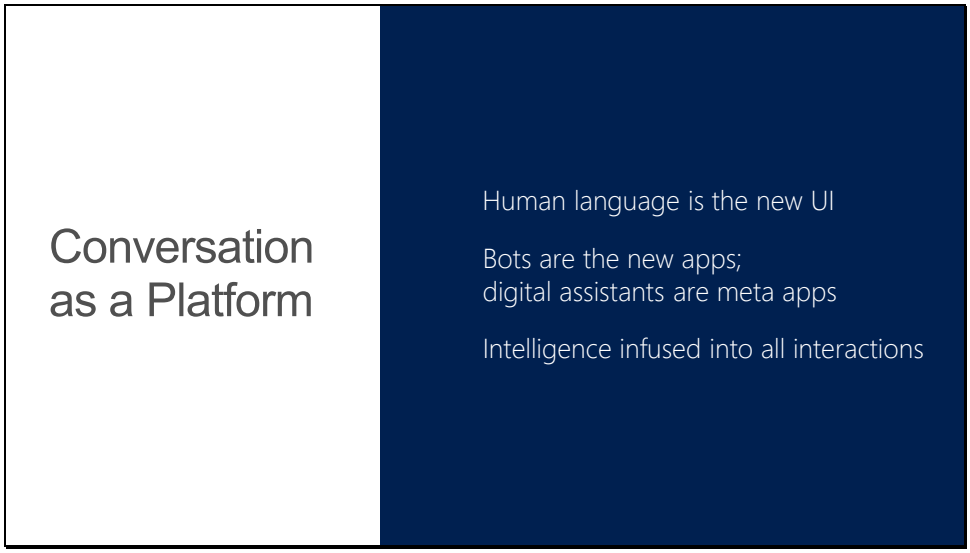
Conversation
as a Platform
Human language is the new UI
Bots are the new apps;
digital assistants are meta apps
Intelligence infused into all interactions

Slide 33

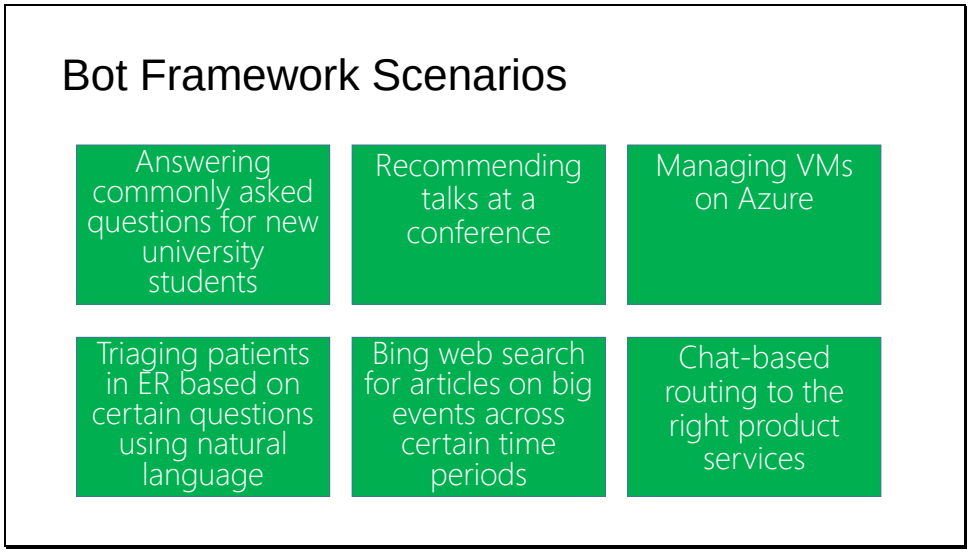
Bot Framework Scenarios
Answering commonly asked questions for new university students
Recommending talks at a conference
Managing VMs on Azure
Triaging patients in ER based on certain questions using natural language
Bing web search for articles on big events across certain time periods
Chat-based routing to the right product services

Slide 34

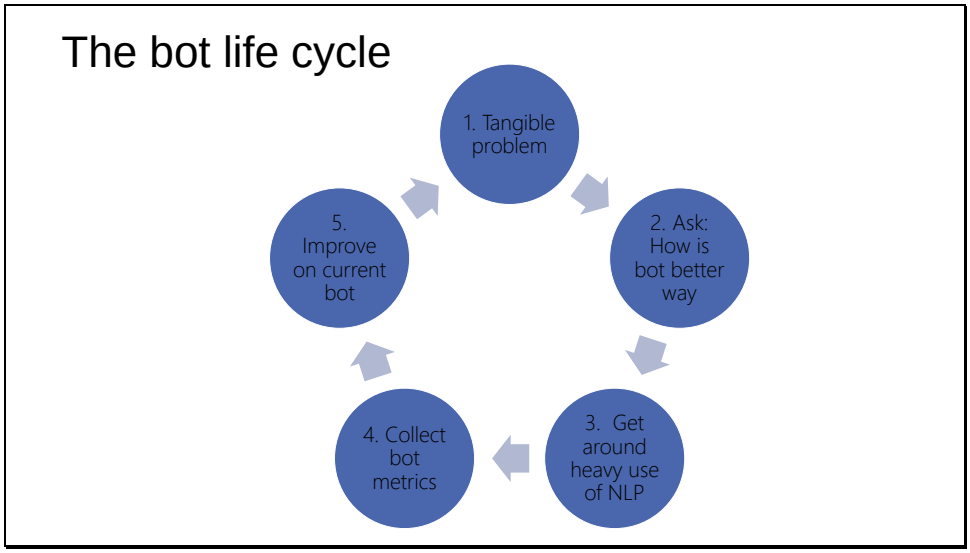

- Start by asking what problem are we trying to solve. Refine until it looks like a tangible problem and not "magic"
- Ask how a bot will be a better experience. User experience is EVERYTHING
- Avoid too much natural language. Careful with unrealistic expectations. Natural language recognition is limited. Menus work great. Commands work great. Buttons, etc.
- You can only analyze and improve your bot if you're collecting metrics for it
- Iterate, improve

Slide 35

Notes on designing an experience users will adopt

Slide 36

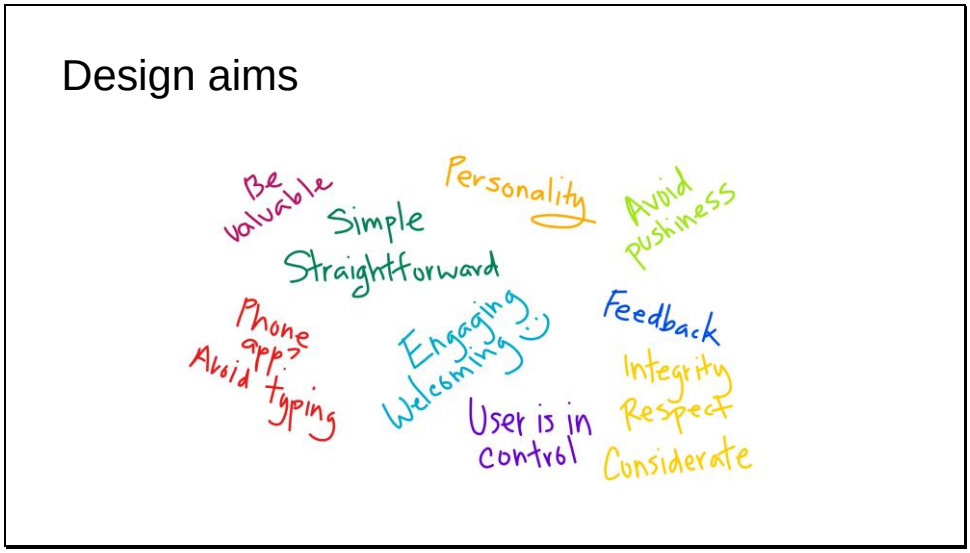

From: https://docs.botframework.com/en-us/directory/best-practices

This may have been in your mind before this tutorial. Ethical and societal considerations as well in an article by Satya Nadella: https://www.linkedin.com/pulse/partnership-future-how-humans-ai-can-work-together-solve-nadella

Slide 37

Creating an awesome experience: the AI

Key considerations around AI and design today:

Outcomes are determined as much by the human element as by the software element.

What does this mean?

The quality of the user experience determines both the usefulness of the product and its rate of adoption, and this is why [it is believed] design is the next frontier of AI.

Credit for paraphrased quotes Manoj Saxena: https://www.fastcodesign.com/3068005/whats-still-missing-from-the-ai-revolution

Slide 38

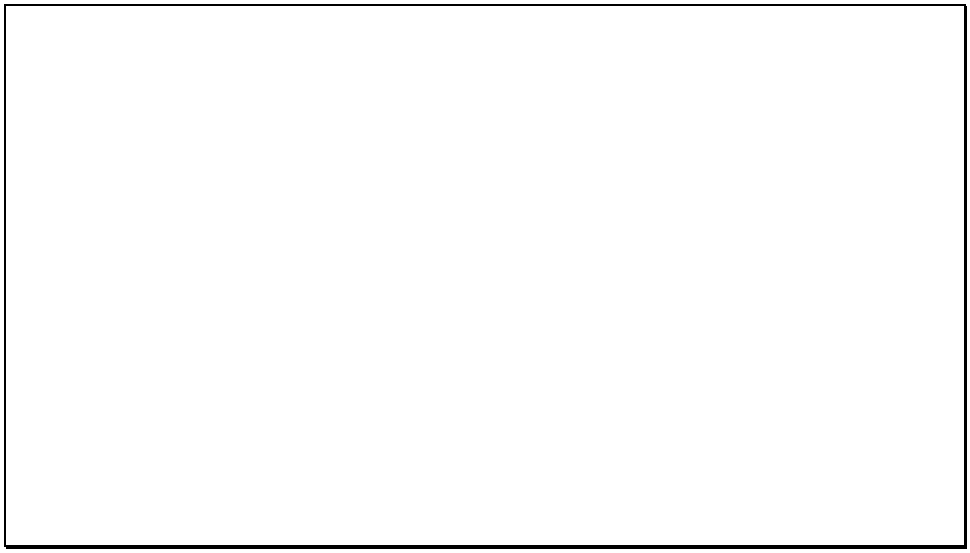

Slide 39

What is the Bot Framework

A development tool and much more…

Slide 40

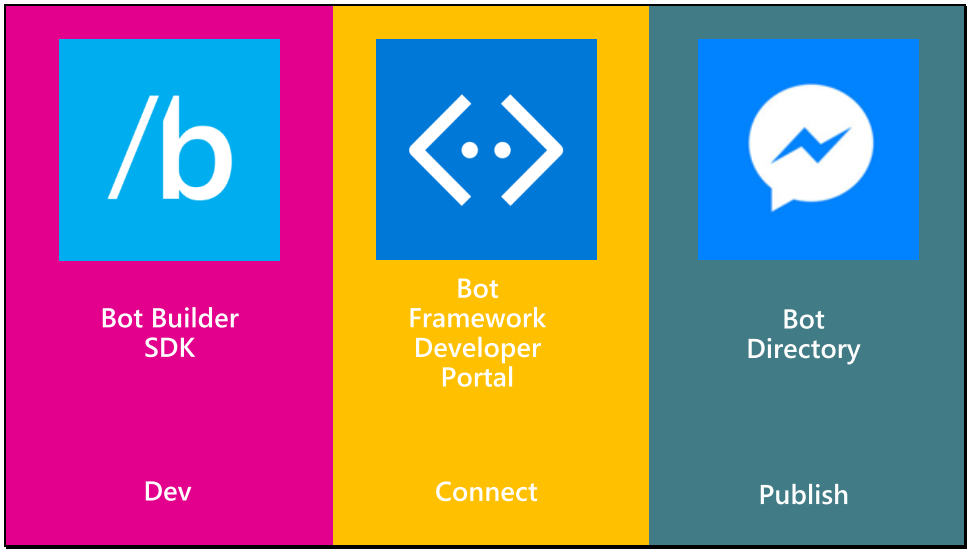

Bot Builder is itself a framework for building conversational applications (“Bots”).

The Bot Builder SDK is an open source SDK hosted on GitHub that provides everything you need to build great dialogs within your Node.js-, .NET- or REST API-based bot.

The Bot Framework Developer Portal lets you connect your bot(s) seamlessly text/sms to Skype, Slack, Facebook Messenger, Kik, Office 365 mail and other popular services. Register, configure and publish.

The Bot Directory is a public directory of all reviewed bots registered through the Developer Portal.

Slide 41

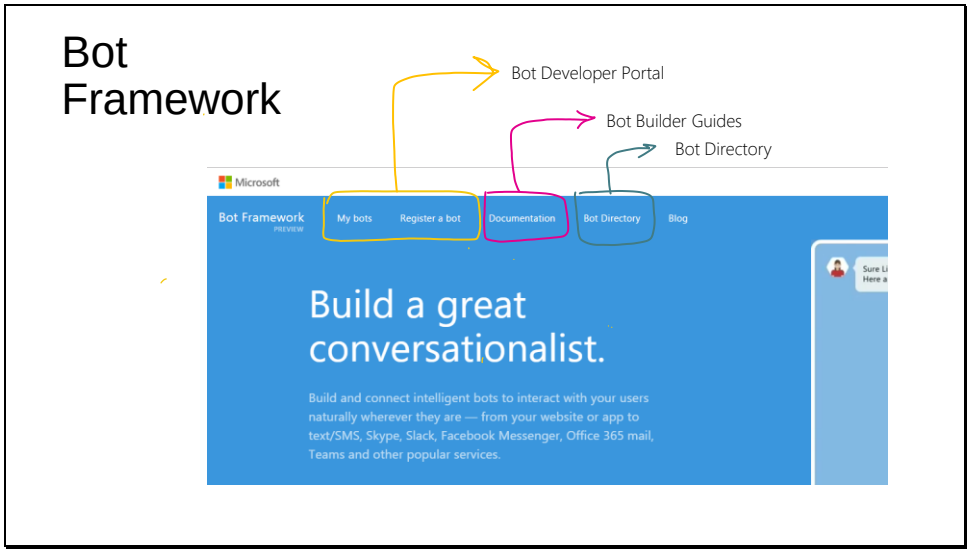

Main page: https://botframework.com

Slide 42

What is the Bot Framework

It’s also a messaging platform

Let’s get a picture of the way it works

Slide 43

What is the Bot Framework: as a development solution

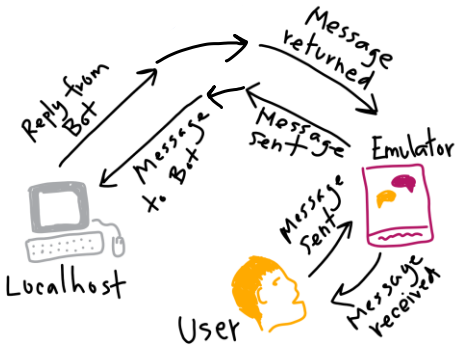

Slide 44

What is the Bot Framework: as a cloud solution

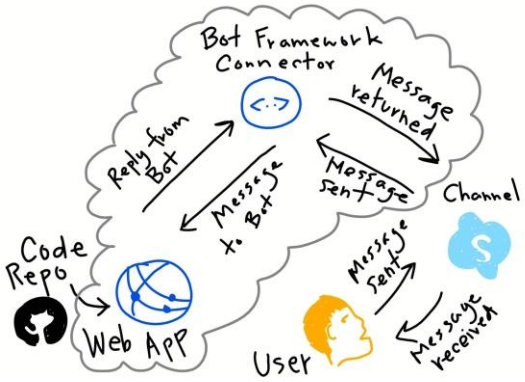

Slide 45

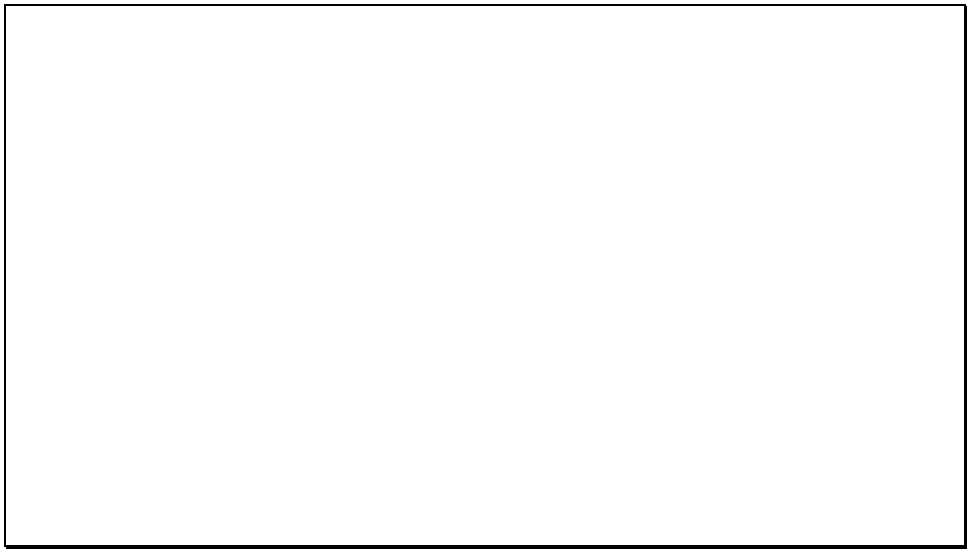

Slide 46

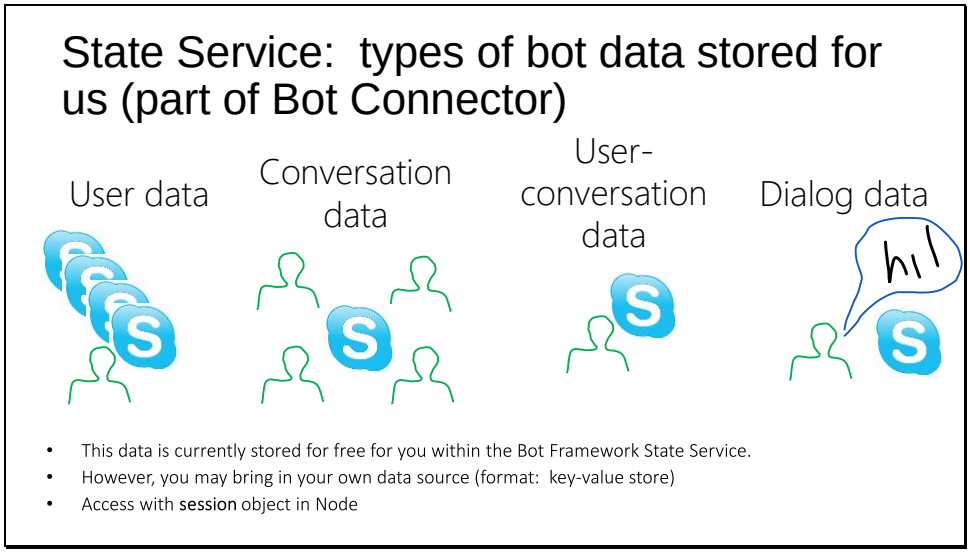

User data – globally available for user across all conversations

conversation data – stores globally for a single conversation (many users could be involved)

User-conversation data – stores globally conversation data for a user (But private to just that user)

Dialog data as well – persists for a single dialog (helpful for temp data in a waterfall set of steps)

If I have a bot that plays Blackjack with me, my stats would be stored in user data (would follow me around from game to game), the deck information and stats in the conversation data (i.e. other players could use the same deck), and my hand in a game would be in user-conversation data (my immediate game's data).

The dialog data persistence ensures that the dialogs state is properly maintained between each turn of the conversation. Dialog data also ensures that a conversation can be picked up later and even on a different machine.

Anything can be stored in these data stores or bags, however it should be limited to data types that are serializable like primitives.

Slide 47

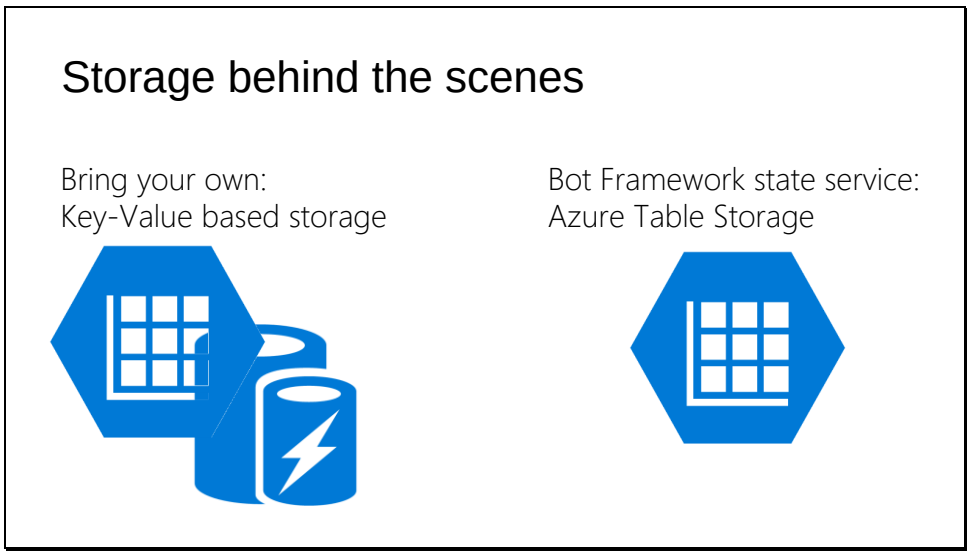

The State REST API has wrappers built around Azure Table Storage. NB, you can bring your own storage in the form of a key-value store like Redis Cache, Table Storage etc. The Bot Framework manages this default storage for you so you can maintain a stateless bot experience and if you bring your own, you'll need to maintain that store and make sure it scales.

Adding your own state example
https://github.com/Microsoft/BotBuilder-Samples/tree/master/Node/core-CustomState

Slide 48

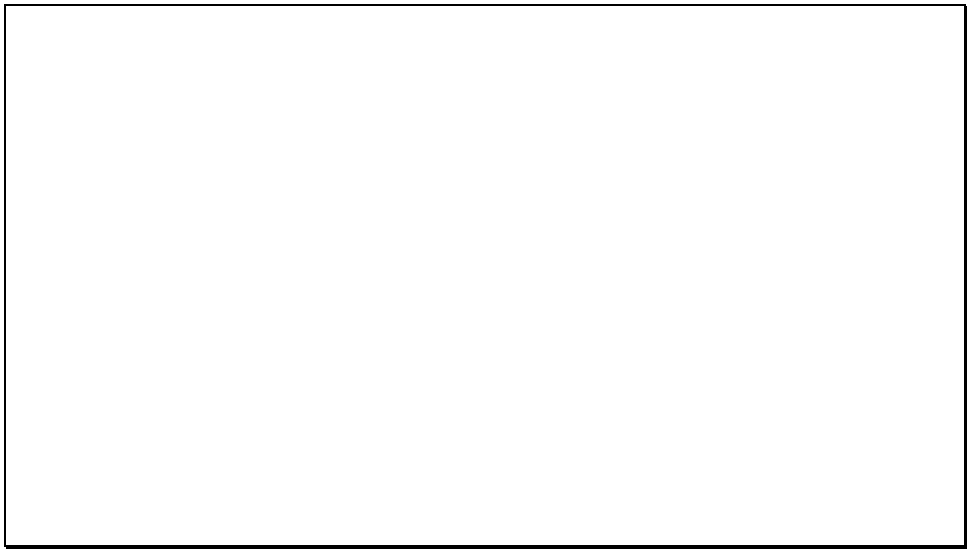

Slide 49

Developing and Deploying Intelligent Chat Bots

Q and A